



Aligning AI with Public Values: Deliberation and Decision-Making for Governing Multimodal LLMs in Political Video Analysis

Truc Bui, Tobias Hösli


Video

2020 Debate Trump vs Biden – On the topic of Climate Change



Choices of Description

Collaborative output brought to you by:



Option 1

- During the climate segment, Trump claims the Paris Accord would cost “tens of millions of jobs” and depicts wind as expensive, intermittent, bird-killing, and dirtier than gas (**wind’s lifecycle CO₂ is far lower than gas [~10–15 vs ~450–500 g/kWh]; turbine bird deaths are orders of magnitude below cats/buildings [FWS.org]).**

For his part, Joe Biden calls climate change an “existential threat to humanity,” **aligning with [IPCC] risk language**, and adds there is a “moral obligation” to act (**normative claim, not a factual assertion**).

Option 2

- Trump speaks fast and **forcefully**, repeating points about jobs, Paris Accord costs, and wind’s downsides over natural gas; **interruptions and emphatic gestures project dominance**. Biden’s delivery is measured; **he looks down, then smiles or laughs at barbs**, framing climate as an “existential threat” and “moral obligation.”


The exchange feels adversarial: Trump drives the tempo and occupies airtime, while Biden counters more with stance and affect than detail, signaling disagreement and values without deep factual sparring.

Research Question

- RQ1: How does the general public perceive the use of LLM in political content interpretation?
 - What do people want from an AI that analyzes a political video?
- RQ2: How do governance/voting mechanisms influence public opinions about improving LLM design?

Methods

Methods



Methods

	Weighted Voting	Quadratic Voting
Equal	100 tokens = 100 votes <i>All users same budget.</i>	2 votes = 4 tokens, 3 votes = 9 tokens, etc. <i>All users same budget.</i>
20/80 Split	100 tokens = 100 votes <i>20% users have 80% budget.</i>	2 votes = 4 tokens, 3 votes = 9 tokens <i>20% users have 80% budget.</i>

Voting Options:

1. Keep current model
- 2. Provide more specific facts**
3. Integrate user feedback loop
4. Analyze speakers' emotions and sentiment.

Results

Divergent Priorities

Experts

- Focused on **emotion, tone, narrative, and context**
- Criticized the AI's summary for lacking crucial **human elements**

Public users

- Preferred **objectivity, clarity, and factual accuracy**
- Wanted fact-checking and **neutral tone**
- Most participants chose **option 2 ("Provide more specific facts")** in all 4 treatment groups.

Results

- Power concentration matters: In the 20/80 condition, the users with more power influenced the final outcome.
- Republicans and Democrats have slightly different attitudes towards the voting mechanisms:
 - Democrats prioritized factual accuracy (option 2).
 - Republicans favored feedback-based improvements (option 3).

	Choice 1		Choice 2		Choice 3		Choice 4	
	mean	std.	mean	std	mean	std	mean	std.
Quadratic - same (n: 29)	0.0814	0.0885	0.4300	0.3421	0.1524	0.1690	0.2597	0.2905
Quadratic -20/80 (n: 30)	0.1193	0.2197	0.3267	0.2260	0.2188	0.1956	0.3080	0.2473
Ranked - same (n: 27)	0.1193	0.1412	0.3941	0.2255	0.1896	0.1197	0.2926	0.2351
Ranked - 20/80 (n:28)	0.1395	0.1534	0.4044	0.2692	0.1877	0.1400	0.2417	0.1870

Table 5: Governance Decision-Making Experience Across Different Political Leaning

Reference to Class Material

- Focuses more on the blockchain implementation and AI governance.
- Computer Vision / Multimodal LLMs: Video→Text pipelines;
 - adding emotion/interaction signals vs factual extraction and fact-checking.

Limitations

- Small N per treatment; results are preliminary. (Which is shown in large standard deviation)
- Group nesting (chat rooms) suggests multilevel/hierarchical standard errors
 - Simple single participant OLS analysis done for Statistical Significance
- Sample skew: more Democrats (51%) & White participants (53%); US-only, single debate clip.
- Counterfactual overreach vs 1-person-1-vote
 - Study only compares QV to weighted voting, but then claims:
 - “in our study, marginalized voices impacted AI design decisions significantly more than under a traditional one person-one-vote system.” which is unsubstantiated

Questions in Plenum

- The research found a clear split in priorities: experts wanted nuance and narrative, while the public wanted facts and objectivity. When aligning an AI, whose preferences should be prioritized?
- Should room expertise change vote weights? If political scientists get 2× votes today, do we like that rule in general?
- If an AI model's design is heavily influenced by a specific political demographic, how do we ensure that the resulting tool is perceived as legitimate and trustworthy across the entire political spectrum, and not just another source of partisan information?

Thanks for Listening

