

Computational Social Science with Images and Audio

Elliott Ash, **Philine Widmer**

20 October 2023

Let us briefly discuss computer vision and (mis)information.

(Inspired by recent sad events around the world.)

- ▶ (Social) Media is flooded with (visual) information on recent tragedies in world politics
 - ▶ Lying anywhere between fully verifiable to entirely fake
- ▶ Verification takes resources (in particular, time)
 - ▶ For one verified or debunked information piece, multiple new ones are posted
- ▶ Many actors, many motives
 - ▶ Are (social) media platforms interested in truth? Citizens? Politicians?
 - ▶ Other motivations than truth?

How can computer vision help/debunk false info?

- ▶ First question: is the image as such fake, or is the meta-information (who, where, what) incorrect?
 - ▶ For example: deep fake vs. wrong context
- ▶ Some helpful approaches:
 - ▶ Reverse image search
 - ▶ Temporal consistency (image metadata)
 - ▶ Image forensics
 - ▶ Face detection
 - ▶ Deepfake detection
 - ▶ Geolocation verification
- ▶ Some concerns
 - ▶ Sophisticated manipulations
 - ▶ Privacy concerns

Let's move back to today's main topic: CNN.

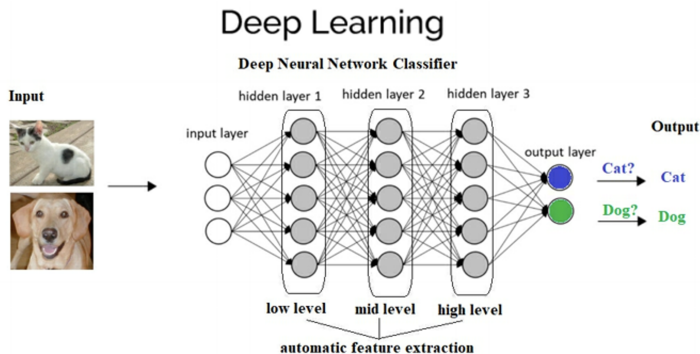


Figure: Dey (2018)¹

¹Dey, S. (2018). Hands-On Image Processing with Python: Expert Techniques for Advanced Image Analysis and Effective Interpretation of Image Data. Packt Publishing Ltd.

Convolutional layers use convolution to filter input data.

- ▶ With an input matrix I and a filter matrix K , the convolution is:

$$(I * K)(x, y) = \sum_{i=-\infty}^{+\infty} \sum_{j=-\infty}^{+\infty} I(i, j) \times K(x - i, y - j)$$

- ▶ Consider this example image:

$$I = \begin{bmatrix} 0 & 2 & 50 \\ 55 & 97 & 230 \\ 234 & 235 & 0 \end{bmatrix}$$

- ▶ ... and a 3×3 filter:

$$K = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$$

What is the dimension of convolved output?

- ▶ What is the convolved output from the previous slide?
- ▶ What dimensions would the output have if the filter was of dimension 2×2 ? Why?
 - ▶ How many times can we put this filter on I above?

$$F = \begin{bmatrix} a & b \\ c & d \end{bmatrix}$$

- ▶ Hint1: It depends on whether we pad the image (e.g., add some zeros at the borders) and...
- ▶ Hint2: ... on the stride, which defines the step size at which the filter evaluates the next position

Convolution is the basis of hierarchical feature learning.

- ▶ Each filter (like K from earlier) produces a unique feature map

$$\text{Image} \xrightarrow{\text{Filter } K} \text{Feature Map}$$

- ▶ A CNN layer typically contains multiple filters
- ▶ Accordingly, an input can produce multiple feature maps in one convolutional layer
 - ▶ In a given layer, if we put filter F_1 and F_2 of 2×2 each (with stride 1 and no padding) on our image I , there will be two feature maps of 2×2 each
 - ▶ Hence, the layer will pass two 2×2 feature maps to the next layer
 - ▶ Disclaimer: heavily simplified for intuition
- ▶ Over training, the network learns the optimal filter values that minimize the loss
- ▶ Pooling layers after convolution can down-sample the feature maps

The above intuition generalizes to multi-channel images.

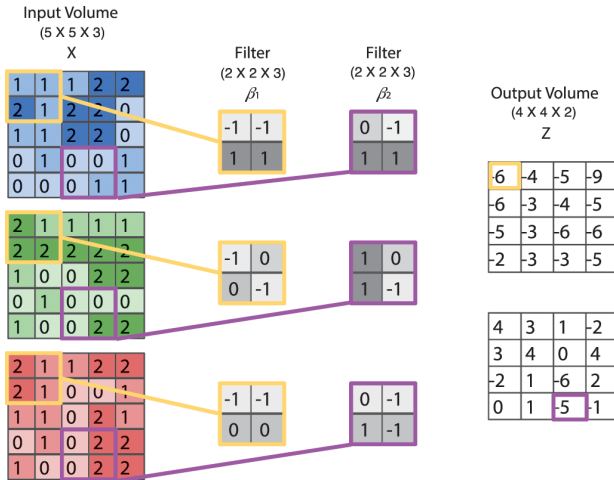


Figure: Convolutional Layer in a CNN for Image Processing (from Webb Williams et al. (2020; cf. first lecture)

CNN come with many hyperparameters.

- ▶ **Epochs/iterations:** Number of times the model trains over the full set of images (*e.g., 100 epochs*)
- ▶ **Learning rate:** How much weights/coefficients change in optimization: essentially scales the gradient (*e.g., 0.01*)
- ▶ **Dropout rate:** To avoid overfitting; share of weights set to 0 (*e.g., 0.5*)
- ▶ **Batch size:** Number of samples processed before updating (*e.g., 32 samples*)
- ▶ **Ratio of train-validation:** Split of images into sets (*e.g., 80-20 split*)
- ▶ **Loss function:** Evaluates model accuracy (*e.g., Mean Squared Error*)

CNN come with many hyperparameters (continued).

- ▶ **Activation functions:** Nonlinear transformations (cf. last lecture) (*e.g., ReLU*)
- ▶ **Optimizer:** Minimizes the loss function (*e.g., Adam*)
- ▶ **Momentum:** Considers previous gradients to avoid stagnation or oscillations (*e.g., 0.9*)
- ▶ **Step size:** Defines learning rate decay interval (*e.g., 10 epochs or 10 batches/iterations*)
- ▶ **Gamma:** Learning rate decay after every step-size (*e.g., 0.1 decay rate*)